

Binary Analysis - 1

REVERSE ENGINEERING

João Paulo Barraca

deti universidade de aveiro
departamento de eletrónica,
telecomunicações e informática

Binary Objects

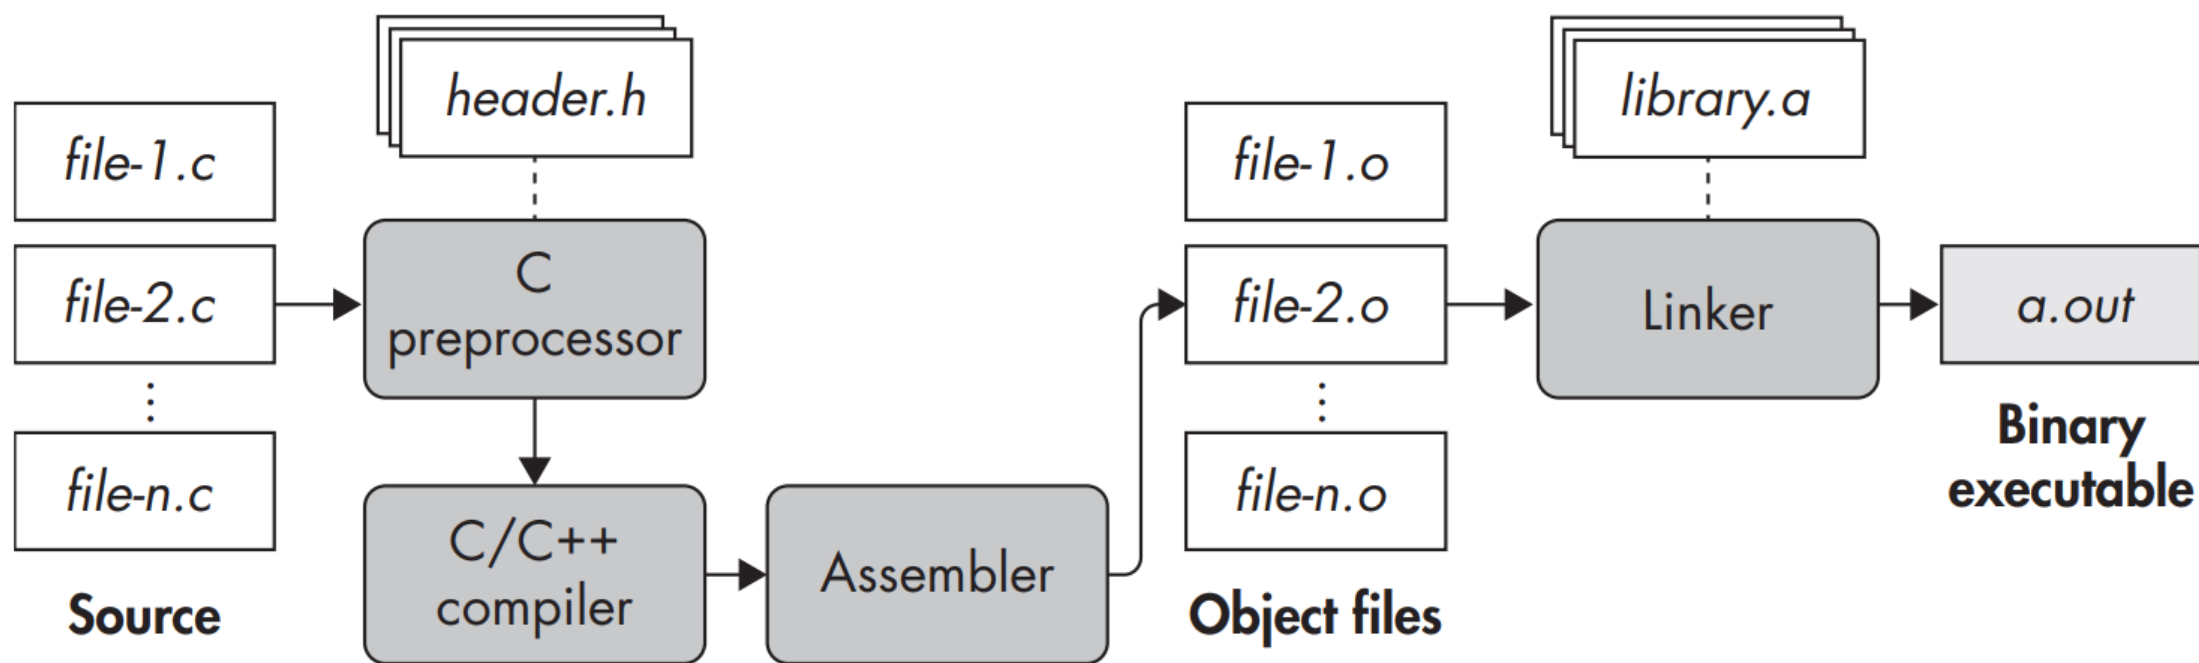


Binary files

- The result of a compilation process
 - Translating high level code (C/C++, etc...) into native code or bytecode
- Code is encapsulated in a binary format
 - It's not a raw file with unstructured bytes
- Target system (CPU or VM) will process the resulting code
 - Which may be only part of the file content

Compilation process

The C/C++ use case

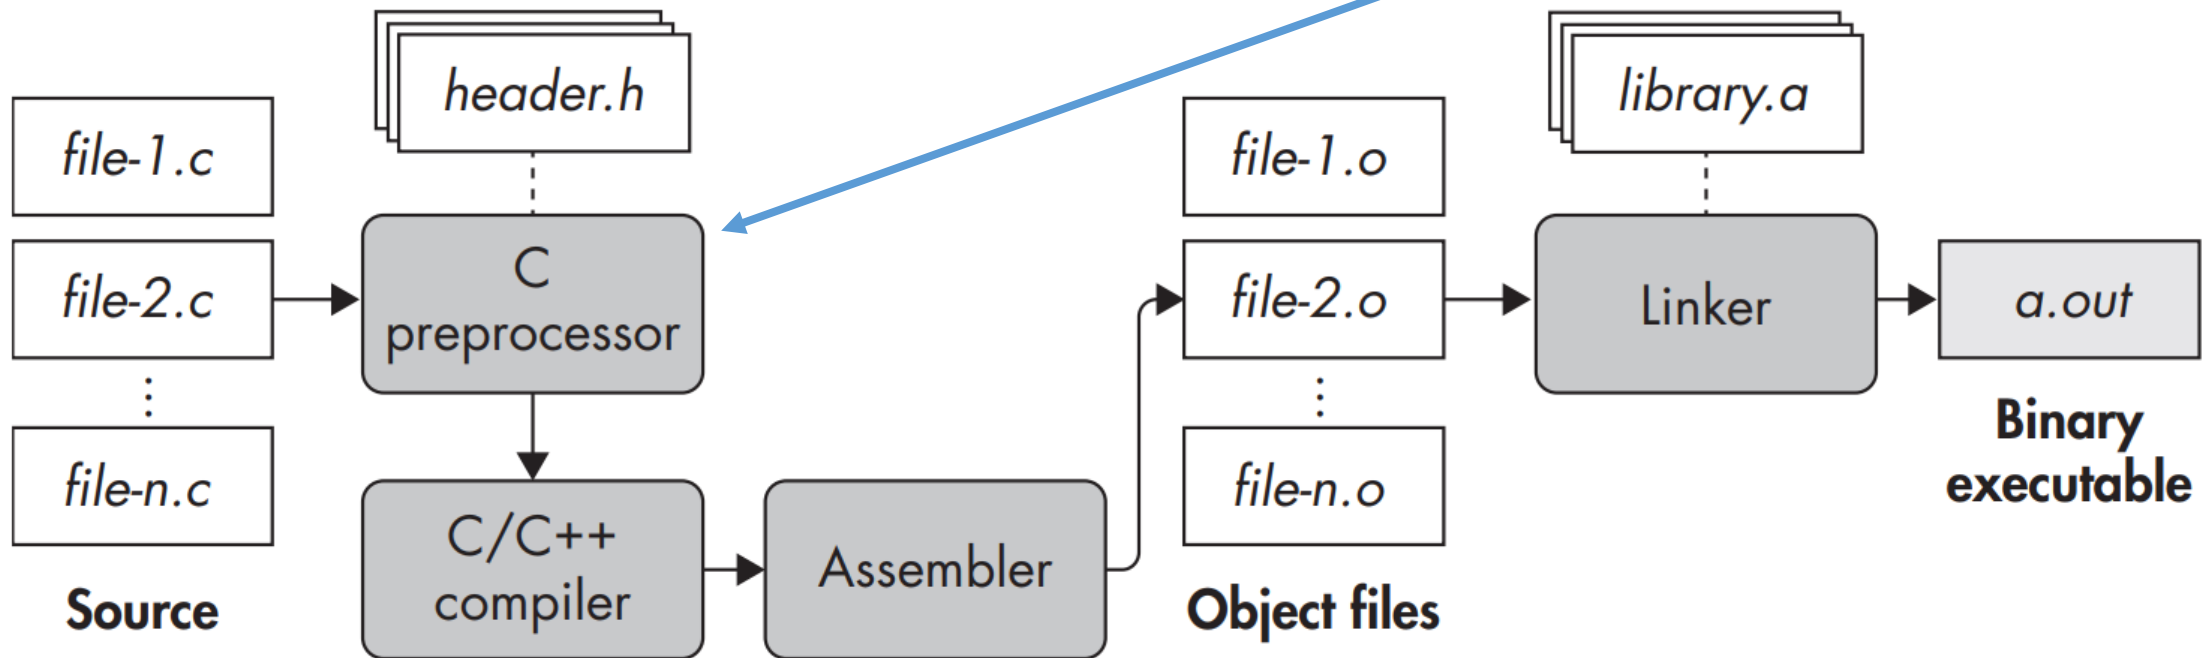


Compilation process

The C/C++ use case

Pre-processor (may be the compiler) processes code, validating its structure and expanding existing macros.

Result is a text blob with content ready to be further processed, and frequently without external dependencies



Hello.c

```
1  #include <stdio.h>
2  #include <stdlib.h>
3
4
5  int main(int argc, char** argv) {
6      printf("Hello World\n");
7
8      return 0;
9  }
10
```

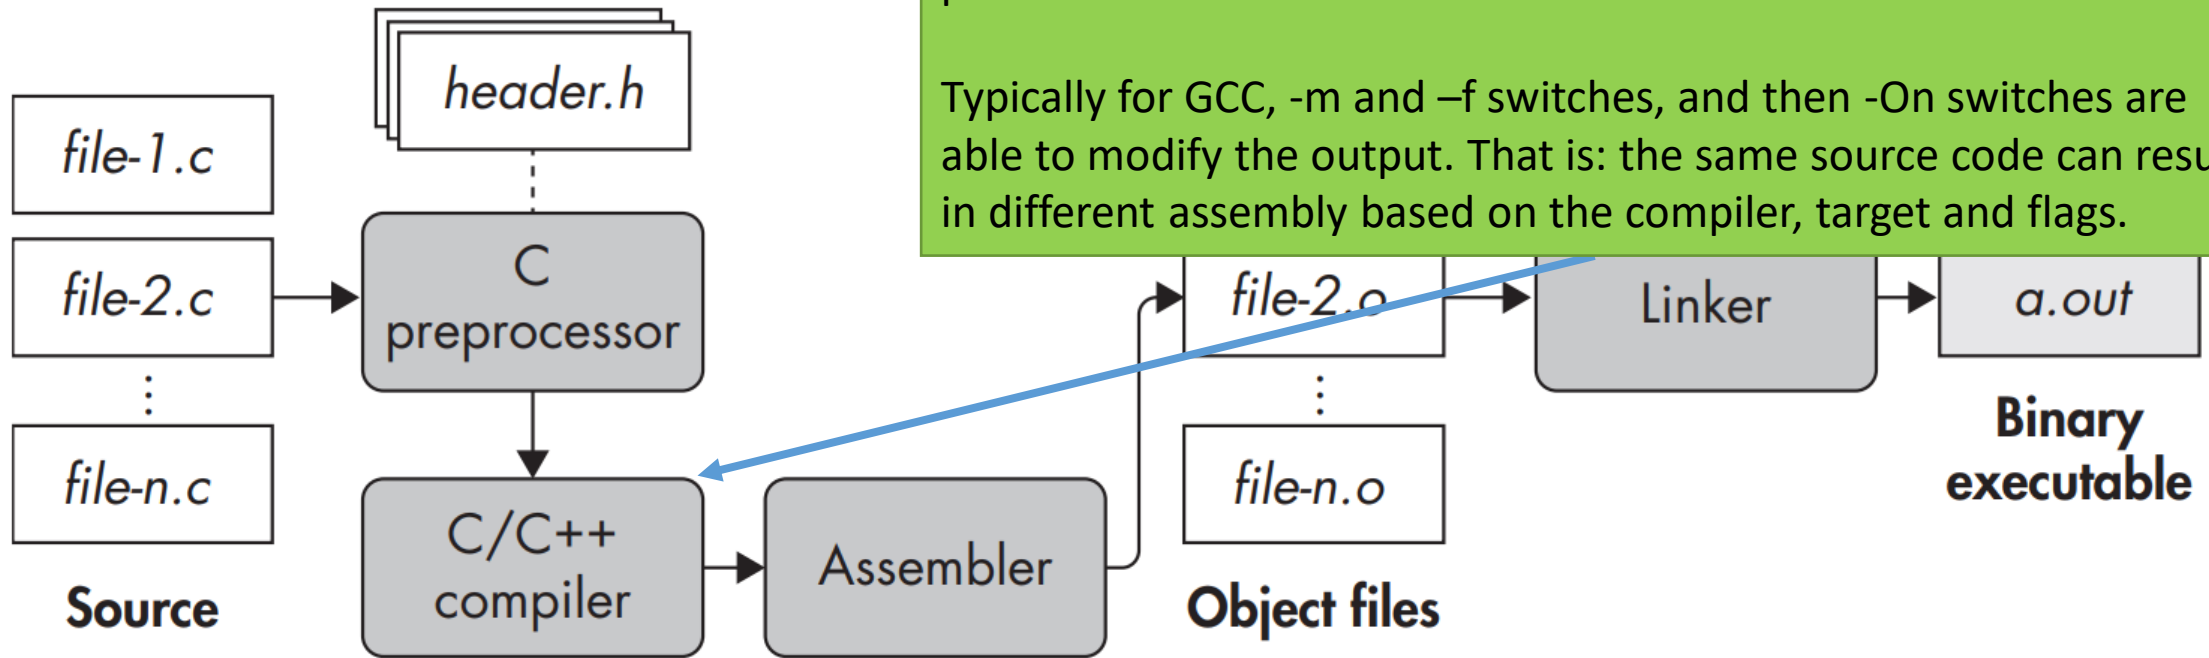
Hello.c

Pre-compile: gcc -E -o hello.e hello.c
produces >1500 lines

```
1760 extern int rpmatch (const char *__response) __attribute__ ((__nothrow__ , __leaf__)) __attribute__ ((__nonnull__ (1))) ;
1761 # 954 "/usr/include/stdlib.h" 3 4
1762 extern int getsubopt (char **__restrict __optionp,
1763     char *const *__restrict __tokens,
1764     char **__restrict __valuep)
1765     __attribute__ ((__nothrow__ , __leaf__)) __attribute__ ((__nonnull__ (1, 2, 3))) ;
1766 # 1000 "/usr/include/stdlib.h" 3 4
1767 extern int getloadavg (double __loadavg[], int __nelem)
1768     __attribute__ ((__nothrow__ , __leaf__)) __attribute__ ((__nonnull__ (1)));
1769 # 1010 "/usr/include/stdlib.h" 3 4
1770 # 1 "/usr/include/x86_64-linux-gnu/bits/stdlib-float.h" 1 3 4
1771 # 1011 "/usr/include/stdlib.h" 2 3 4
1772 # 1020 "/usr/include/stdlib.h" 3 4
1773
1774 # 3 "hello.c" 2
1775
1776
1777
1778 # 5 "hello.c"
1779 int main(int argc, char** argv) {
1780     printf("Hello World\n");
1781
1782     return 0;
1783 }
```

Compilation process

The C/C++ use case



Compiler processes the file and produces assembly code. This may result in assembly for an intermediate processor, and not the final processor.

The processor will create abstract syntax trees (AST) and may tweak or optimize the result according to the options it was provided with.

Typically for GCC, -m and -f switches, and then -O switches are able to modify the output. That is: the same source code can result in different assembly based on the compiler, target and flags.

Hello.c

Compile: `gcc -masm intel -S -o hello.s hello.c`

File Metadata

Constant variables and symbols

Compiler additional data. In this case Call Frame Information to handle exceptions

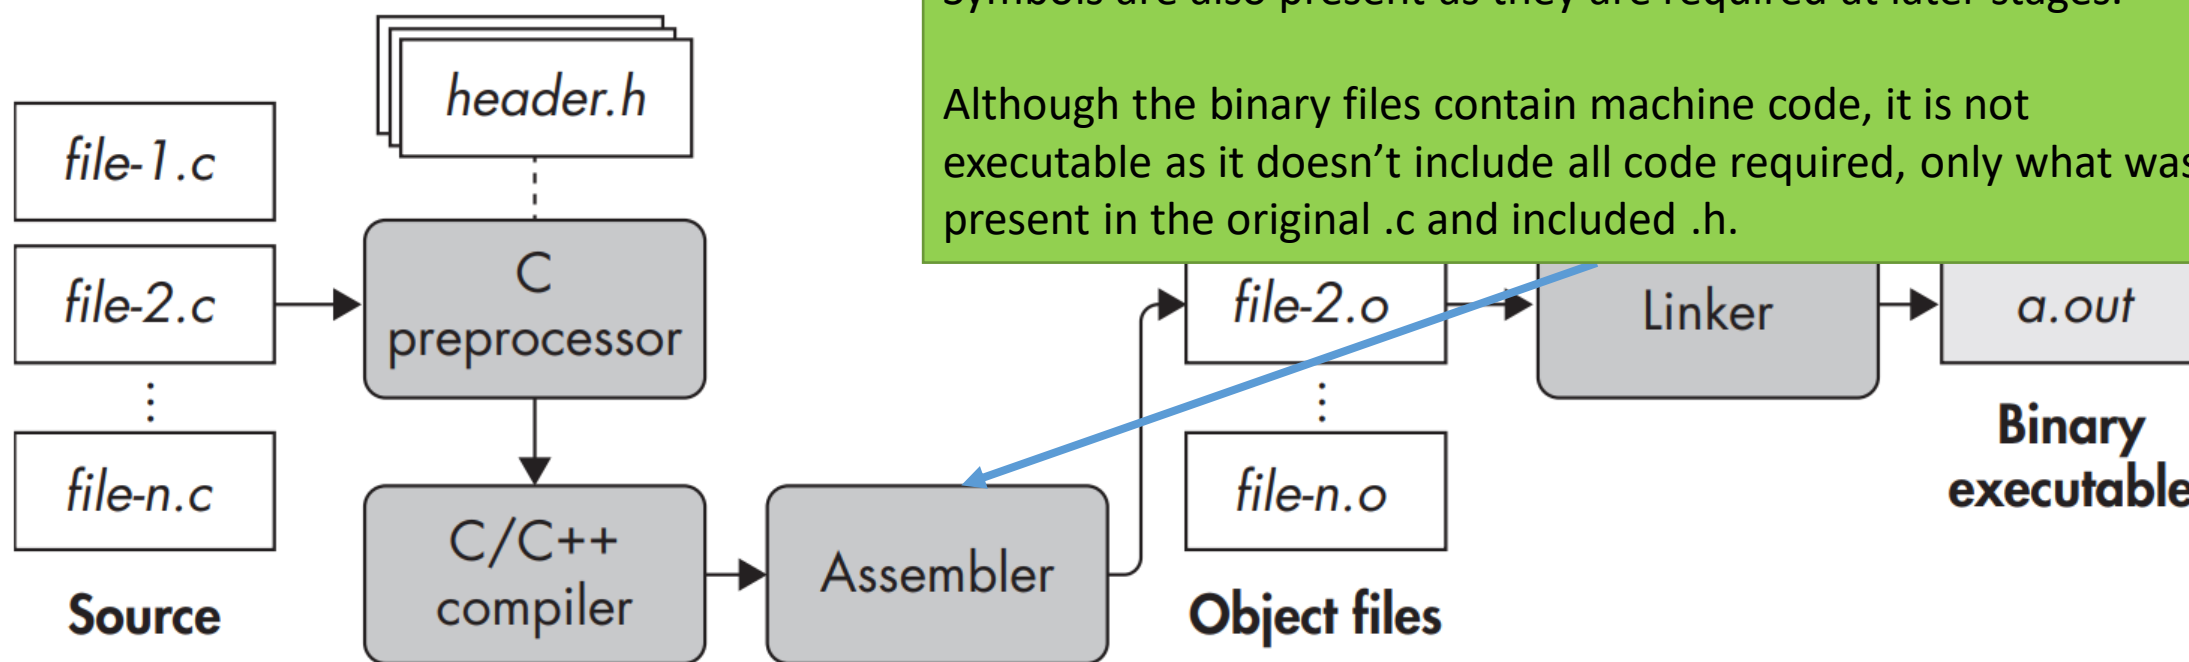
Assembly instructions. Notice that symbols are kept as labels

Additional sections to produce:
Entry point
Compiler identification
Instruct linker to mark stack as NX

```
1  .file "hello.c"
2  .intel_syntax noprefix
3  .text
4  .section .rodata
5  .LC0:
6  .string "Hello World"
7  .text
8  .globl main
9  .type main, @function
10 main:
11 .LFB6:
12 .cfi_startproc
13 push rbp
14 .cfi_def_cfa_offset 16
15 .cfi_offset 6, -16
16 mov rbp, rsp
17 .cfi_def_cfa_register 6
18 sub rsp, 16
19 mov DWORD PTR -4[rbp], edi
20 mov QWORD PTR -16[rbp], rsi
21 lea rdi, .LC0[rip]
22 call puts@PLT
23 mov eax, 0
24 leave
25 .cfi_def_cfa 7, 8
26 ret
27 .cfi_endproc
28 .LFE6:
29 .size main, .-main
30 .ident "GCC: (Debian 8.3.0-6) 8.3.0"
31 .section .note.GNU-stack,"",@progbits
```

Compilation process

The C/C++ use case



Input containing assembly code is transformed into machine code. Output is a set of object files, or modules with a `.o` extension.

Code produced may use relative addresses, making it reusable (technically *relocatable*) when integrated into a final binary file.

Symbols are also present as they are required at later stages.

Although the binary files contain machine code, it is not executable as it doesn't include all code required, only what was present in the original `.c` and included `.h`.

Hello.c

Compile: `gcc -c -o hello.o hello.c`

Assemble the code into machine code

File is an Executable and Linkable Format (ELF)

Cannot be executed

Defines a symbol **main** in the Text section

_GLOBAL_OFFSET_TABLE_ and **puts** are not defined.

Code is not present on the object file

64 bit, Least Significant Byte (Little Endian)

Not stripped = Contains symbols

```
$ gcc -c -o hello.o hello.c

$ file hello.o
hello.o: ELF 64-bit LSB relocatable, x86-64, version 1 (SYSV), not stripped

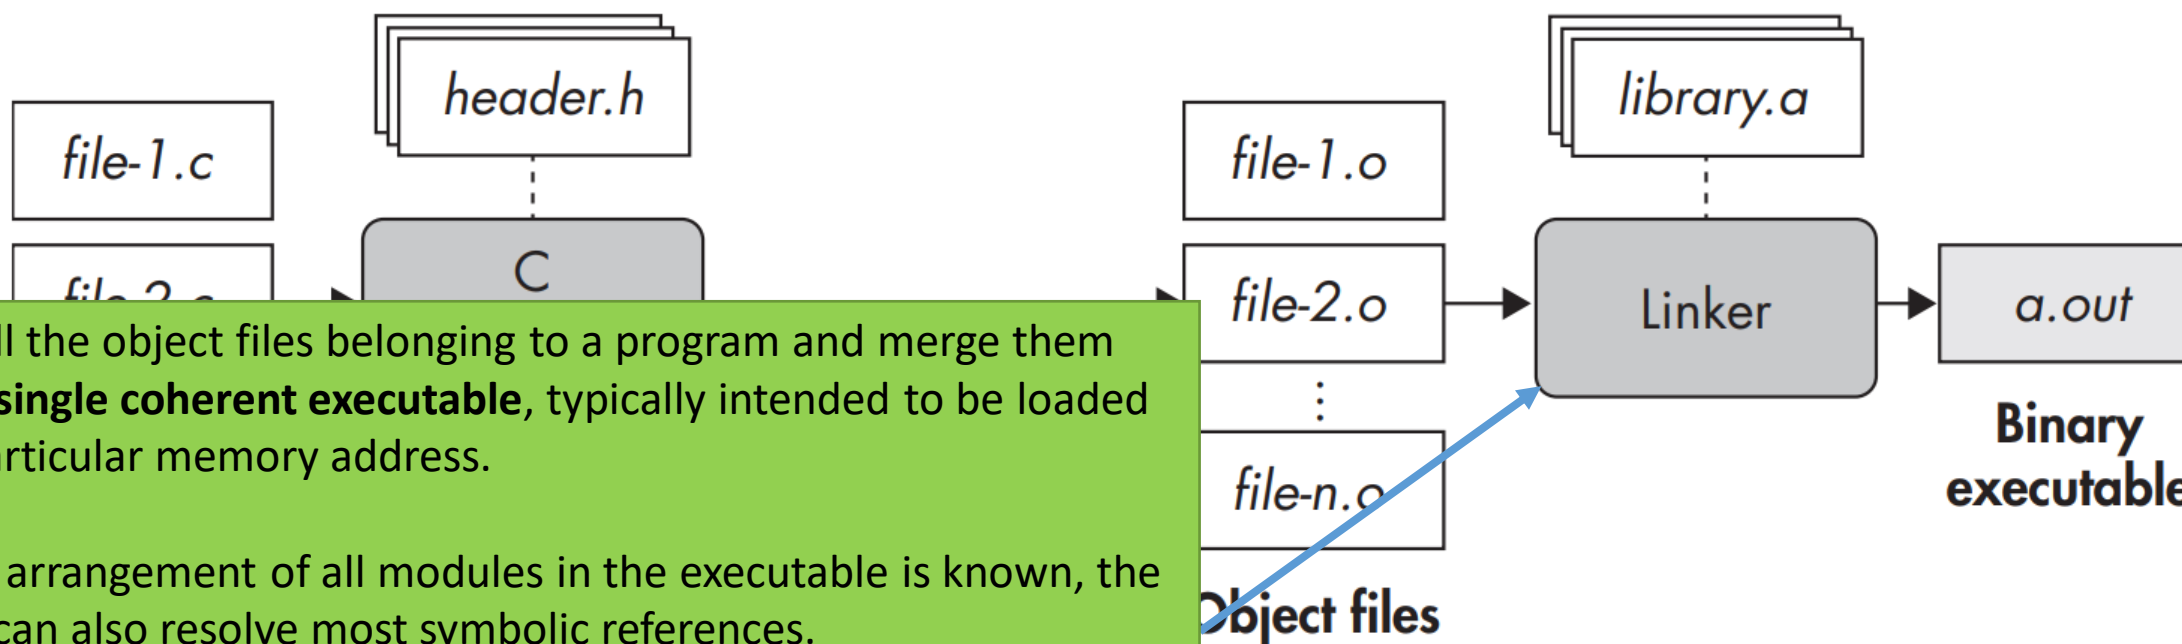
$ ./hello.o
-bash: ./hello.o: cannot execute binary file: Exec format error

$ nm hello.o
0000000000000000 U _GLOBAL_OFFSET_TABLE_
                 T main
                 U puts
```

SYSV = System V ABI. Identifies the target system (others: Solaris, Tru64, FreeBSD, NetBSD...)

Compilation process

The C/C++ use case



Take all the object files belonging to a program and merge them into a **single coherent executable**, typically intended to be loaded at a particular memory address.

As the arrangement of all modules in the executable is known, the linker can also resolve most symbolic references.

References to libraries may or may not be completely resolved, depending on the type of library. In this case, the library is added as a dependency and the symbol is resolved in real time.

Hello.c

Compile: `gcc -o hello hello.c`

- 64 bit, Little Endian Architecture
- Position Independent Executable (Can use ALSR)
- Uses shared libraries
- Uses the ld-linux-x86-64.so.2 loader
- sha1 build id
- Not stripped: contains symbol names

```
1 $ gcc -o hello hello.c
2
3 $ file hello
4 hello: ELF 64-bit LSB pie executable, x86-64, version 1 (SYSV), dynamically linked, interpreter
5 /lib64/ld-linux-x86-64.so.2, for GNU/Linux 3.2.0, BuildID[sha1]=2d3c19e9d0110eef7554245eb02d70bcc9b60dd2, not stripped
6
7 $ ldd hello
8      linux-vdso.so.1 (0x00007ffffdb8a1000)
9      libc.so.6 => /lib/x86_64-linux-gnu/libc.so.6 (0x00007f3e844a0000)
10     /lib64/ld-linux-x86-64.so.2 (0x00007f3e84685000)
```

Shared libraries required to execute this file. Some code is not on the hello binary and is on the libraries

Hello.c

```
$ nm hello
```

```
0000000000004030 B __bss_start
0000000000004030 b completed.7325
               w __cxa_finalize@@GLIBC_2.2.5
0000000000004020 D __data_start
0000000000004020 W data_start
0000000000001080 t deregister_tm_clones
00000000000010f0 t __do_global_dtors_aux
0000000000003df0 t __do_global_dtors_aux_fini_array_entry
0000000000004028 D __dso_handle
0000000000003df8 d _DYNAMIC
0000000000004030 D _edata
0000000000004038 B _end
00000000000011c4 T _fini
0000000000001130 t frame_dummy
0000000000003de8 t __frame_dummy_init_array_entry
0000000000002154 r __FRAME_END__
0000000000004000 d _GLOBAL_OFFSET_TABLE__
               w __gmon_start__
0000000000002010 r __GNU_EH_FRAME_HDR
0000000000001000 t _init
0000000000003df0 t __init_array_end
0000000000003de8 t __init_array_start
0000000000002000 R _IO_stdin_used
               w _ITM_deregisterTMCloneTable
               w _ITM_registerTMCloneTable
00000000000011c0 T __libc_csu_fini
0000000000001160 T __libc_csu_init
               U __libc_start_main@@GLIBC_2.2.5
0000000000001135 T main
               U puts@@GLIBC_2.2.5
00000000000010b0 t register_tm_clones
0000000000001050 T _start
0000000000004030 D __TMC_END__
```

Symbols present in the file

Bb: in the BSS

D: in the initialized data Sec.

Rr: in the Read Only Data Sec.

Tt: in the Text (code) Sec.

U: Undefined

Ww: Weak

- default impl. to be overridden

Some are undefined.
Will be defined by the dynamic linker. Code resides on an external object.

Executable Symbols

Tables

- Symbols are **names identifying addresses of a binary**
 - Have a type, such as Function, and including Undefined
 - E.g. functions create symbols, especially external functions (puts)
- ELF files have **two symbol tables**
 - **.dynsym**: symbols which will be allocated to memory when the program loads.
 - In the example, puts is provided by libc, required for operation, and exists as a dynamic symbol
 - **.symtab**: contains all symbols, including many used for linking and debugging, but not related to code required for execution.
 - These areas will not be allocated (mapped) to RAM
 - Extremely useful to identify the name of functions/sections when reversing!

Executable Symbols

Stripping

- Only symbols in the .dyntab are required
 - Identify allocated sections
 - Identify symbols that must be resolved in external libraries
 - Used for Dynamic Linking when the program is loaded
- **Stripping** is the process of removing unused symbols and code from a binary
 - Stripped binaries take less space, and are not reversed so easily
 - There is no hints about the purpose of a function from its name

Hello

```
$ readelf --syms hello
```

```
Symbol table '.dynsym' contains 7 entries:
```

Num:	Value	Size	Type	Bind	Vis	Ndx	Name
0:	0000000000000000	0	NOTYPE	LOCAL	DEFAULT	UND	
1:	0000000000000000	0	NOTYPE	WEAK	DEFAULT	UND	__ITM_deregisterTMCloneTab
2:	0000000000000000	0	FUNC	GLOBAL	DEFAULT	UND	puts@GLIBC_2.2.5 (2)
3:	0000000000000000	0	FUNC	GLOBAL	DEFAULT	UND	__libc_start_main@GLIBC_2.2.5 (2)
4:	0000000000000000	0	NOTYPE	WEAK	DEFAULT	UND	__gmon_start__
5:	0000000000000000	0	NOTYPE	WEAK	DEFAULT	UND	__ITM_registerTMCloneTable
6:	0000000000000000	0	FUNC	WEAK	DEFAULT	UND	__cxa_finalize@GLIBC_2.2.5 (2)

```
Symbol table '.symtab' contains 64 entries:
```

Num:	Value	Size	Type	Bind	Vis	Ndx	Name
0:	0000000000000000	0	NOTYPE	LOCAL	DEFAULT	UND	
1:	00000000000002a8	0	SECTION	LOCAL	DEFAULT	1	
2:	00000000000002c4	0	SECTION	LOCAL	DEFAULT	2	
...							
48:	0000000000000000	0	FUNC	GLOBAL	DEFAULT	UND	puts@@GLIBC_2.2.5
49:	0000000000004030	0	NOTYPE	GLOBAL	DEFAULT	24	__edata
50:	00000000000011c4	0	FUNC	GLOBAL	HIDDEN	15	__fini
51:	0000000000000000	0	FUNC	GLOBAL	DEFAULT	UND	__libc_start_main@@GLIBC_
52:	0000000000004020	0	NOTYPE	GLOBAL	DEFAULT	24	__data_start
53:	0000000000000000	0	NOTYPE	WEAK	DEFAULT	UND	__gmon_start__
54:	0000000000004028	0	OBJECT	GLOBAL	HIDDEN	24	__dso_handle
55:	0000000000002000	4	OBJECT	GLOBAL	DEFAULT	16	__IO_stdin_used
56:	0000000000001160	93	FUNC	GLOBAL	DEFAULT	14	__libc_csu_init
57:	0000000000004038	0	NOTYPE	GLOBAL	DEFAULT	25	__end
58:	0000000000001050	43	FUNC	GLOBAL	DEFAULT	14	__start
59:	0000000000004030	0	NOTYPE	GLOBAL	DEFAULT	25	__bss_start
60:	0000000000001135	34	FUNC	GLOBAL	DEFAULT	14	main
61:	0000000000004030	0	OBJECT	GLOBAL	HIDDEN	24	__TMC_END__
62:	0000000000000000	0	NOTYPE	WEAK	DEFAULT	UND	__ITM_registerTMCloneTable
63:	0000000000000000	0	FUNC	WEAK	DEFAULT	UND	__cxa_finalize@@GLIBC_2.2

Hello.c

Binary is stripped of extra symbols

Only the .dynsym table is kept
Required for identifying allocatable areas
Notice as all symbols here are undefined (must be dynamically linked)

```
1 $ strip hello
2
3 $ readelf --syms hello
4
5 Symbol table '.dynsym' contains 7 entries:
6   Num:      Value              Size Type    Bind   Vis      Ndx Name
7   0: 0000000000000000          0 NOTYPE  LOCAL  DEFAULT UND
8   1: 0000000000000000          0 NOTYPE  WEAK   DEFAULT UND _ITM_deregisterTMCloneTab
9   2: 0000000000000000          0 FUNC    GLOBAL DEFAULT UND puts@GLIBC_2.2.5 (2)
10  3: 0000000000000000          0 FUNC    GLOBAL DEFAULT UND __libc_start_main@GLIBC_2.2.5 (2)
11  4: 0000000000000000          0 NOTYPE  WEAK   DEFAULT UND __gmon_start__
12  5: 0000000000000000          0 NOTYPE  WEAK   DEFAULT UND _ITM_registerTMCloneTable
13  6: 0000000000000000          0 FUNC    WEAK   DEFAULT UND __cxa_finalize@GLIBC_2.2.5 (2)
```

What is inside an Object File?

- An Object File contains information required to execute a program (not only code)
 - May not include all implementation, as this can be dynamically loaded
- Information is kept in sections, which are processed differently. Some are:
 - **.rodata**: readonly data, containing strings
 - **.got**: Global Offset Table - maps symbols to memory locations (offsets).
 - **.plt**: **Procedure Linkage Table** – uses the PLT to transfer execution to the correct location of a symbol, dealing with external symbols and fixing the GOT
 - **.bss**: **Block Starting Symbol** – contains uninitialized variables
 - **.dynsym**: List of symbols in allocatable memory
 - ... many others:
 - To read sections: `readelf -S hello`
 - To dump all code: `objdump -M intel -d hello`

Hello ELF content

RODATA: objdump -sj .rodata

Contains Read Only Data (Strings and other constants)

```
1  $ objdump -sj .rodata hello
2
3  hello:      file format elf64-x86-64
4
5  Contents of section .rodata:
6  2000 01000200 48656c6c 6f20576f 726c6400 ....Hello World.
```

Hello ELF disassembly

Indirection at PLT

The entry point to the program.

Prepares stack

Calls **main** function

The main function:

Allocates 0x10 in the stack

Sets arguments to puts

Calls puts@PLT

Sets the Return Code to 0

Leave

```
1  objdump -M intel -d hello
2
3  hello:      file format elf64-x86-64
4
5  ...
6
7  0000000000001030 <puts@plt>:
8      1030:      ff 25 e2 2f 00 00      jmp     QWORD PTR [rip+0x2fe2]          # 4018 <puts@GLIBC_2.2.5>
9      1036:      68 00 00 00 00      push    0x0
10     103b:      e9 e0 ff ff ff      jmp     1020 <.>
11
12  ...
13
14  0000000000001050 <_start>:
15     1050:      31 ed      xor     ebp,ebp
16     1052:      49 89 d1    mov     r9,rdx
17     1055:      5e      pop     rsi
18     1056:      48 89 e2    mov     rdx,rsi
19     1059:      48 83 e4 f0 and     rsp,0xfffffffffffffff0
20     105d:      50      push    rax
21     105e:      54      push    rsp
22     105f:      4c 8d 05 5a 01 00 00 lea     r8,[rip+0x15a]          # 11c0 <__libc_csu_fini>
23     1066:      48 8d 0d f3 00 00 00 lea     rcx,[rip+0xf3]          # 1160 <__libc_csu_init>
24     106d:      48 8d 3d c1 00 00 00 lea     rdi,[rip+0xc1]          # 1135 <main>
25     1074:      ff 15 66 2f 00 00    call    QWORD PTR [rip+0x2f66]      # 3fe0 <__libc_start_main@GLIBC_2.2.5>
26     107a:      f4      hlt
27     107b:      0f 1f 44 00 00      nop     DWORD PTR [rax+rax*1+0x0]
28
29  0000000000001135 <main>:
30     1135:      55      push    rbp
31     1136:      48 89 e5    mov     rbp,rsi
32     1139:      48 83 ec 10 sub     rsp,0x10
33     113d:      89 7d fc    mov     DWORD PTR [rbp-0x4],edi
34     1140:      48 89 75 f0 mov     QWORD PTR [rbp-0x10],rsi
35     1144:      48 8d 3d b9 0e 00 00 lea     rdi,[rip+0xeb9]          # 2004 <_IO_stdin_used+0x4>
36     114b:      e8 e0 fe ff ff      call    1030 <puts@plt>
37     1150:      b8 00 00 00 00      mov     eax,0x0
38     1155:      c9      leave
39     1156:      c3      ret
40     1157:      66 0f 1f 84 00 00 00 nop     WORD PTR [rax+rax*1+0x0]
41     115e:      00 00
```

Hello Relocations

`readelf --relocs hello`

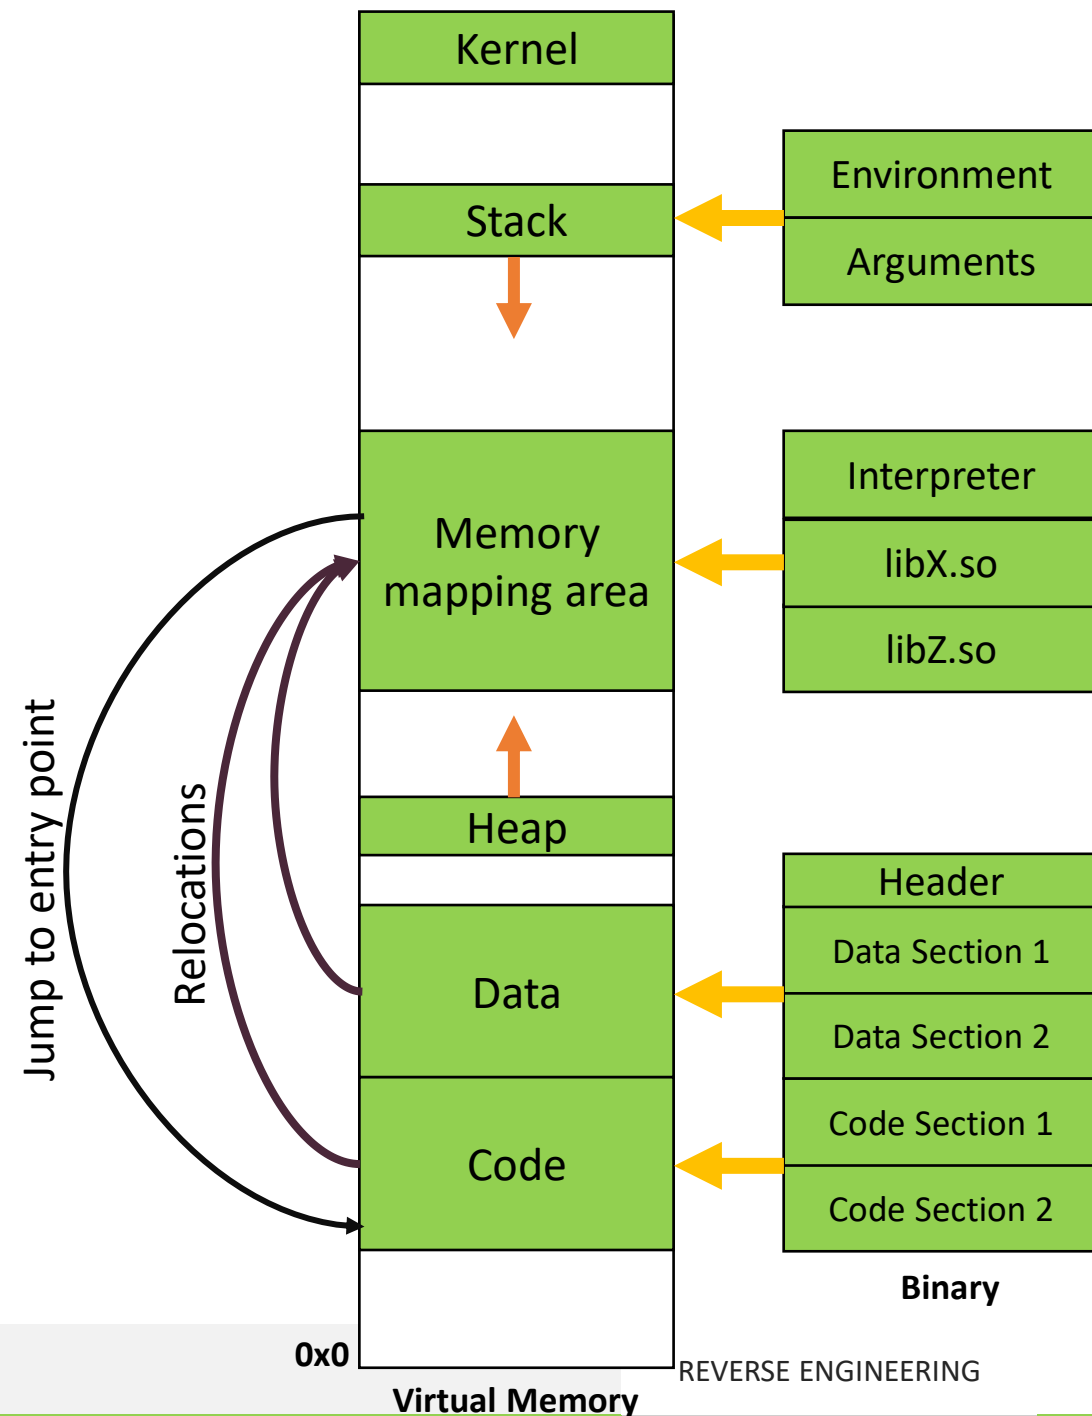
```
1 $ readelf --relocs hello
2
3 ...
4
5 Relocation section '.rela.plt' at offset 0x548 contains 1 entry:
6   Offset          Info          Type           Sym. Value      Sym. Name + Addend
7   000000004018    000200000007  R_X86_64_JUMP_SLO 0000000000000000 puts@GLIBC_2.2.5 + 0
```

Symbol to be dynamically linked

LIBC TAG to match

How are objects loaded?

- File is split according to existing sections
 - Each loaded at a different location (with different access attributes)
- Libraries are also mapped in the program address space
 - All code from libraries is present
- Stack grows downwards, heap grows upwards
 - On modern SOs, growth may be limited, not on microcontrollers
- Interpreter is required to setup the binary in memory
 - `ld-Linux.so` or `ntdll.dll`
 - `readelf -p .interp filename`
 - Will handle relocations, resolving required symbols
 - If lazy-loading is used, relocation is done when the symbol is first used

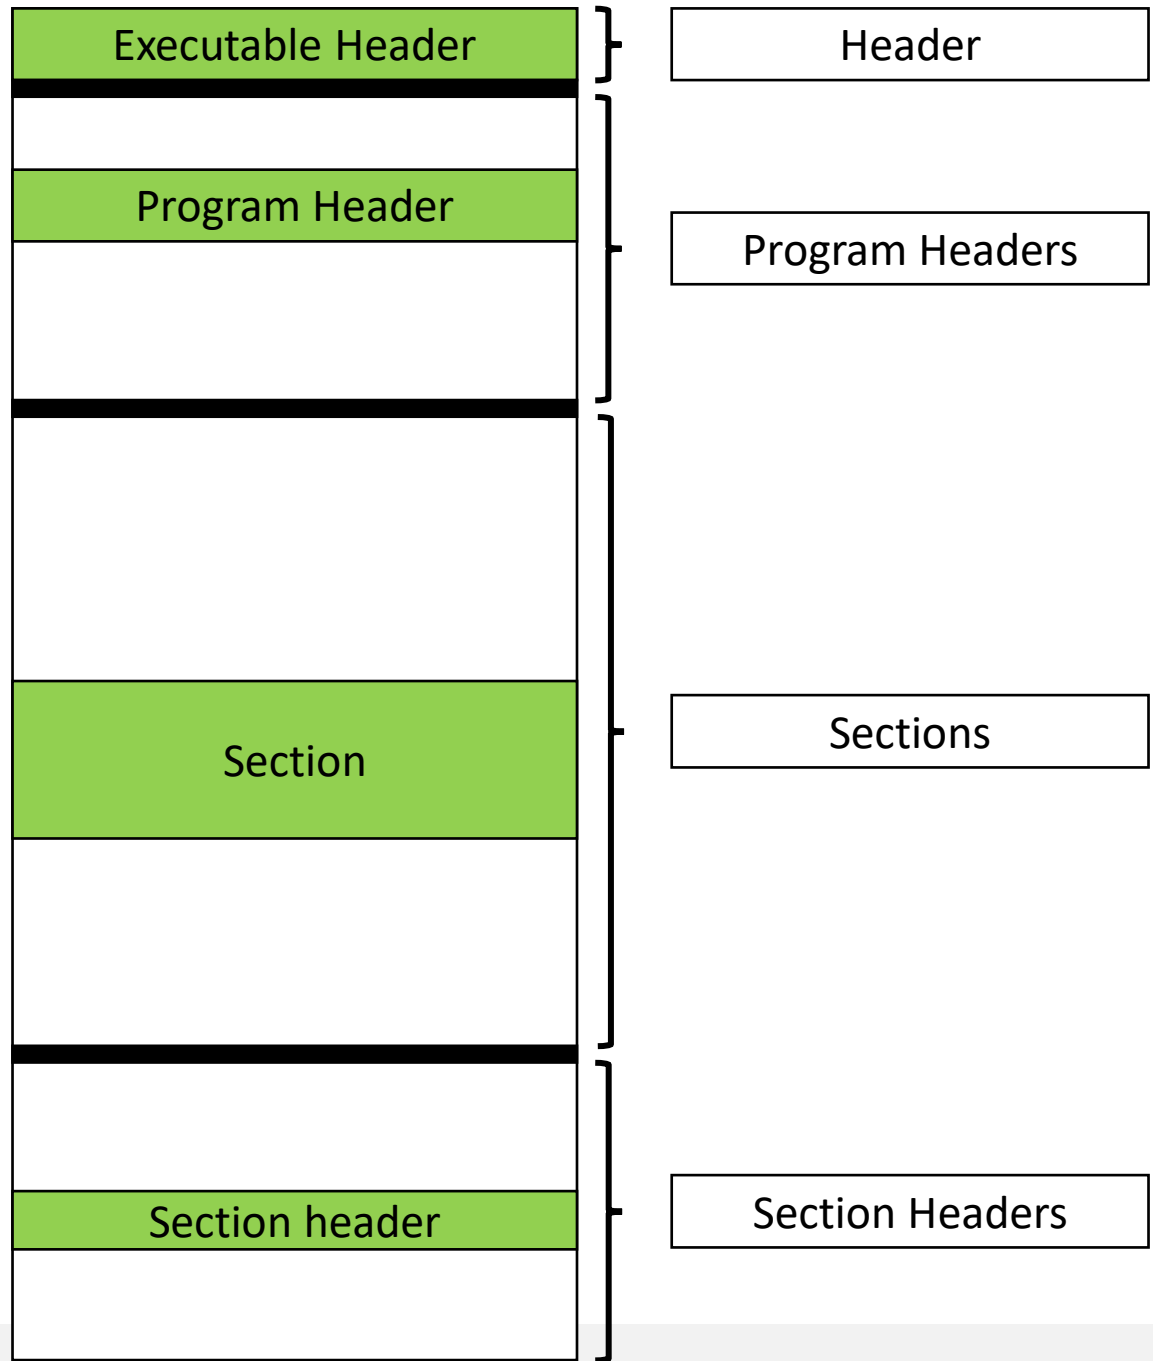


ELF Files



ELF – Executable and Linkable Format

- Container for executable files, object files, shared libraries, and core dumps
 - And other things out of this context like in Android
- Composed by several headers and sections:
 - Executable Header
 - Several Program Headers (optional)
 - Several Sections, with a header and content



ELF Headers

Executable Header

- Mandatory header, with basic information about the file
 - Architecture
 - Entry Point
 - Header locations and number
 - Type
 - Type of data
- Follow the structure **Elf64_Ehdr**
 - defined in `/usr/include/elf.h`

```
1 $ readelf -h hello
2
3 ELF Header:
4   Magic:   7f 45 4c 46 02 01 01 00 00 00 00 00 00 00 00 00
5   Class:                               ELF64
6   Data:                                   2's complement, little endian
7   Version:                             1 (current)
8   OS/ABI:                               UNIX - System V
9   ABI Version:                          0
10  Type:                                  DYN (Shared object file)
11  Machine:                               Advanced Micro Devices X86-64
12  Version:                               0x1
13  Entry point address:                   0x1050
14  Start of program headers:              64 (bytes into file)
15  Start of section headers:              14688 (bytes into file)
16  Flags:                                  0x0
17  Size of this header:                    64 (bytes)
18  Size of program headers:                56 (bytes)
19  Number of program headers:              11
20  Size of section headers:                64 (bytes)
21  Number of section headers:              30
22  Section header string table index: 29
```

ELF Headers

Section Headers

- Sections are unstructured placeholders of data (frequently code) **targeting the Linker**
 - Some sections are well known and follow a defined structure
 - Some sections can be arbitrary binary blob
 - Some sections may contain content not useful for execution
 - Section order is irrelevant
 - Symbols, relocation information is stored in sections
- Section headers describe the properties of each section
 - Name, type, flags, address when loaded, file offset, size, information...
- ELF files that require no linking, may omit section headers

```
1 $ readelf -S hello |grep "\["
```

	[Nr]	Name	Type	Address	Offset
2	[0]		NULL	0000000000000000	00000000
3	[1]	.interp	PROGBITS	00000000000002a8	000002a8
4	[2]	.note.ABI-tag	NOTE	00000000000002c4	000002c4
5	[3]	.note.gnu.build-i	NOTE	00000000000002e4	000002e4
6	[4]	.gnu.hash	GNU_HASH	0000000000000308	00000308
7	[5]	.dynsym	DYNSYM	0000000000000330	00000330
8	[6]	.dynstr	STRTAB	00000000000003d8	000003d8
9	[7]	.gnu.version	VERSYM	000000000000045a	0000045a
10	[8]	.gnu.version_r	VERNEED	0000000000000468	00000468
11	[9]	.rela.dyn	RELA	0000000000000488	00000488
12	[10]	.rela.plt	RELA	0000000000000548	00000548
13	[11]	.init	PROGBITS	0000000000001000	00001000
14	[12]	.plt	PROGBITS	0000000000001020	00001020
15	[13]	.plt.got	PROGBITS	0000000000001040	00001040
16	[14]	.text	PROGBITS	0000000000001050	00001050
17	[15]	.fini	PROGBITS	00000000000011c4	000011c4
18	[16]	.rodata	PROGBITS	0000000000002000	00002000
19	[17]	.eh_frame_hdr	PROGBITS	0000000000002010	00002010
20	[18]	.eh_frame	PROGBITS	0000000000002050	00002050
21	[19]	.init_array	INIT_ARRAY	0000000000003de8	00002de8
22	[20]	.fini_array	FINI_ARRAY	0000000000003df0	00002df0
23	[21]	.dynamic	DYNAMIC	0000000000003df8	00002df8
24	[22]	.got	PROGBITS	0000000000003fd8	00002fd8
25	[23]	.got.plt	PROGBITS	0000000000004000	00003000
26	[24]	.data	PROGBITS	0000000000004020	00003020
27	[25]	.bss	NOBITS	0000000000004030	00003030
28	[26]	.comment	PROGBITS	0000000000000000	00003030
29	[27]	.symtab	SYMTAB	0000000000000000	00003050
30	[28]	.strtab	STRTAB	0000000000000000	00003650
31	[29]	.shstrtab	STRTAB	0000000000000000	00003853

ELF Sections

.init and .fini

- Contains executable code required before/after the binary entry point is executed
 - Initialization tasks to prepare/clean the memory space
- Some uses:
 - prepare profiling tasks (`__gmon_start__`)
 - Invoke global constructors/destructors (C++)
 - Save program arguments

```
1 $ objdump -M intel -d -j .init hello
2
3 hello:      file format elf64-x86-64
4
5
6 Disassembly of section .init:
7
8 00000000000001000 <_init>:
9      1000:      48 83 ec 08          sub     rsp,0x8
10     1004:      48 8b 05 dd 2f 00 00  mov     rax,QWORD PTR [rip+0x2fdd]      # 3fe8 <__gmon_start__>
11     100b:      48 85 c0            test    rax,rax
12     100e:      74 02             je      1012 <_init+0x12>
13     1010:      ff d0            call    rax
14     1012:      48 83 c4 08          add     rsp,0x8
15     1016:      c3              ret
```

ELF Sections

.text section

- Contains the main program code
 - The main target of a Reverse Engineering activity
 - Allocated as executable and read-only
 - Contains the user code, and additional code created by the compiler
 - Cleanup/initialization functions, stack guards, etc..
- In this section resides the program entry point
 - When the binary is loaded, execution flow is transferred that address
 - **Related** to the **main** function in a C program (but not the main)

ELF Sections

.text section: Entry Point

The **hello** program entry point address

```
1 $ objdump -M intel -d -j .text hello
2
3 hello:      file format elf64-x86-64
4
5
```

Disassembly of section **.text**:

```
8 0000000000001050 <_start>:
9      1050:      31 ed      xor     ebp,ebp
10     1052:      49 89 d1    mov     r9,rdx
11     1055:      5e         pop     rsi
12     1056:      48 89 e2    mov     rdx,rsi
13     1059:      48 83 e4 f0 and     rsp,0xfffffffffffffff0
14     105d:      50         push    rax
15     105e:      54         push    rsp
16     105f:      4c 8d 05 5a 01 00 00 lea     r8,[rip+0x15a]      # 11c0 <__libc_csu_fini>
17     1066:      48 8d 0d f3 00 00 00 lea     rcx,[rip+0xf3]      # 1160 <__libc_csu_init>
18     106d:      48 8d 3d c1 00 00 00 lea     rdi,[rip+0xc1]      # 1135 <main>
19     1074:      ff 15 66 2f 00 00    call   QWORD PTR [rip+0x2f66] # 3fe0 <__libc_start_main@GLIBC_2.2.5>
20     107a:      f4         hlt
21     107b:      0f 1f 44 00 00      nop     DWORD PTR [rax+rax*1+0x0]
22
23 ...
```

Loads the address of the main function into **RDI** (first argument) of a function

Calls [__libc_start_main@GLIBC 2.2.5](#) which transfers control to the program **main** function

ELF Sections

.bss, .data, .rodata

- **.rodata**: Read only data
 - Stores constant values
 - Mapped to a page marked as read only
- **.data**: Area with information to initialize variables
 - As the data can be modified, the section is writable
- **.bss**: Unitialized variables
 - Memory is allocated for a variable that may be required, but nothing else is done
 - As there is no data associated, the .bss doesn't take space on the binary. Only instructs the system to reserve memory.

ELF Sections

.plt, .got, .got.plt

- **Procedure Linkage Table and Global Offset Table**
 - **.PLT**: Code to relocate symbols
 - **.GOT**: Array with addresses of each symbol requiring relocation
 - **.got** is similar to **.got.plt** but it's writable, while **.got** may be marked as Read Only as a security measure (-z relro)
 - Using a table (GOT) allows patching this table, while keeping libraries in same address, shared to multiple processes
- **Sections required for lazy binding (real time relocation)**
 - Linker needs to resolve the effective address of a code identified by a symbol (e.g **puts**)
 - The code may be on the program, or on an external library, mapped to the virtual memory
 - **.plt** and **.got** ensure the symbol location is found and the code jumps around correctly
 - This is executed as the symbols are required! (LAZY)
 - On Linux, the Env Variable **LD_BIND_NOW** forces linking by the linker (on program load)
 - Will increase performance during execution, but will slow down startup

ELF Sections

Lazy Binding

(1) The **puts** function is called. The function is on an external library, and it must be relocated. So, it jumps to the **puts@plt**

```
1 0000000000001020 <.plt>:
2   1020:      push    QWORD PTR [rip+0x2fe2]      # 4008 <_GLOBAL_OFFSET_TABLE_+0x8>
3   1026:      jmp     QWORD PTR [rip+0x2fe4]      # 4010 <_GLOBAL_OFFSET_TABLE_+0x10>
4   102c:      nop     DWORD PTR [rax+0x0]
5
6 0000000000001030 <puts@plt>:
7   1030:      jmp     QWORD PTR [rip+0x2fe2]      # 4018 <puts@GLIBC_2.2.5>
8   1036:      push    0x0
9   103b:      jmp     1020 <.plt>
10
11 ....
12
13 0000000000001135 <main>:
14 ...
15   114b:      call    1030 <puts@plt>
16 ...
17
18 0000000000004000 <_GLOBAL_OFFSET_TABLE_>:
19   4000:      f8 3d 00 00 00 00 00 00 00 00 00 00 00 00 00 00  .=.....
20   ...
21   4018:      36 10 00 00 00 00 00 00                          6.....
```

ELF Sections

Lazy Binding

(2) At the PLT, the code doesn't jump to the final location, as it is not known (yet)

Instead, it jumps to an entry at the GOT (0x4018). In this case, the value is 0x1036, pointing to the code at line 8.

Remember: This is a static analysis, the dynamic linker is not working, so the symbol is unresolved

```
1 0000000000001020 <.plt>:
2   1020:      push    QWORD PTR [rip+0x2fe2]      # 4008 <_GLOBAL_OFFSET_TABLE_+0x8>
3   1026:      jmp     QWORD PTR [rip+0x2fe4]      # 4010 <_GLOBAL_OFFSET_TABLE_+0x10>
4   102c:      nop     DWORD PTR [rax+0x0]
5
6 → 0000000000001030 <puts@plt>:
7   1030:      jmp     QWORD PTR [rip+0x2fe2]      # 4018 <puts@GLIBC_2.2.5>
8   1036:      push    0x0
9   103b:      jmp     1020 <.plt>
10
11 ....
12
13 0000000000001135 <main>:
14 ...
15   114b:      call    1030 <puts@plt>
16 ...
17
18 0000000000004000 <_GLOBAL_OFFSET_TABLE_>:
19   4000:      f8 3d 00 00 00 00 00 00 00 00 00 00 00 00 00 00  .=.
20   ...
21   4018:      36 10 00 00 00 00 00 00                        6.....
```

ELF Sections

Lazy Binding

(3) A value 0 is pushed. This is an identifier that is stored to the stack. An index, actually.

The code then jumps to the .plt generic functions at 0x1020.

A new identifier is pushed (the address in the GOT that is missing the entry)

Code jumps to the Dynamic Linker

```
1 0000000000001020 <.plt>:
2   1020:      push    QWORD PTR [rip+0x2fe2]      # 4008 <_GLOBAL_OFFSET_TABLE_+0x8>
3   1026:      jmp     QWORD PTR [rip+0x2fe4]      # 4010 <_GLOBAL_OFFSET_TABLE_+0x10>
4   102c:      nop     DWORD PTR [rax+0x0]
5
6 → 0000000000001030 <puts@plt>:
7   1030:      jmp     QWORD PTR [rip+0x2fe2]      # 4018 <puts@GLIBC_2.2.5>
8   1036:      push    0x0
9   103b:      jmp     1020 <.plt>
10
11 ....
12
13 0000000000001135 <main>:
14 ...
15   114b:      call    1030 <puts@plt>
16 ...
17
18 0000000000004000 <_GLOBAL_OFFSET_TABLE_>:
19   4000:      f8 3d 00 00 00 00 00 00 00 00 00 00 00 00 00 00  .=.
20   ...
21   4018:      36 10 00 00 00 00 00 00 6.....
```

ELF Sections

Lazy Binding

(4) At the dynamic linker, it searches for the symbols in the mapped libraries and writes a value to the GOT at 0x4018.

Then he calls that address.

```
1 0000000000001020 <.plt>:
2   1020:      push    QWORD PTR [rip+0x2fe2]      # 4008 <_GLOBAL_OFFSET_TABLE_+0x8>
3   1026:      jmp     QWORD PTR [rip+0x2fe4]      # 4010 <_GLOBAL_OFFSET_TABLE_+0x10>
4   102c:      nop     DWORD PTR [rax+0x0]
5
6 0000000000001030 <puts@plt>:
7   1030:      jmp     QWORD PTR [rip+0x2fe2]      # 4018 <puts@GLIBC_2.2.5>
8   1036:      push    0x0
9   103b:      jmp     1020 <.plt>
10
11 ....
12
13 0000000000001135 <main>:
14 ...
15   114b:      call    1030 <puts@plt>
16 ...
17
18 0000000000004000 <_GLOBAL_OFFSET_TABLE_>:
19   4000:      f8 3d 00 00 00 00 00 00 00 00 00 00 00 00 00 00  .=.....
20   ...
21   4018:      36 10 00 00 00 00 00 00                          6.....
```

ELF Sections

Lazy Binding

(2.1) At the PLT, the code doesn't jump to the final location, as it is not known (yet)

Instead, it jumps to an entry at the GOT (**0x4018**).

If the program is executing, and it is the second time puts is called, the entry has **0x7ffffff651910**, which points to the real puts.

This was obtained by loading the binary in GDB and using GEF

```
1 0000000000001020 <.plt>:
2   1020:      push    QWORD PTR [rip+0x2fe2]      # 4008 <_GLOBAL_OFFSET_TABLE_+0x8>
3   1026:      jmp     QWORD PTR [rip+0x2fe4]      # 4010 <_GLOBAL_OFFSET_TABLE_+0x10>
4   102c:      nop     DWORD PTR [rax+0x0]
5
6 → 0000000000001030 <puts@plt>:
7   1030:      jmp     QWORD PTR [rip+0x2fe2]      # 4018 <puts@GLIBC_2.2.5>
8   1036:      push    0x0
9   103b:      jmp     1020 <.plt>
10
11 ....
12
13 0000000000001135 <main>:
14 ...
15   114b:      call    1030 <puts@plt>
16 ...
17
18 gef> got
19
20 GOT protection: Partial RelRO | GOT functions: 1
21
[0x8004018] puts@GLIBC_2.2.5 → 0x7ffffff651910
```

ELF Sections

.rel.*, .rela.*

- Tables containing information to the dynamic linker about the required relocations
 - **R_X86_64_GLOB_DAT**: GOT offset should be filled with the symbol address (Lines 8-12)
 - **R_X86_64_JUMP_SLO**: Jump Slots to be represented in the **.got.plt** and **.plt** sections as shown previously (Line 16)

```
1 $ readelf --relocs hello
2
3 Relocation section '.rela.dyn' at offset 0x488 contains 8 entries:
4   Offset          Info          Type           Sym. Value      Sym. Name + Addend
5   000000003de8     000000000008 R_X86_64_RELATIVE      1130
6   000000003df0     000000000008 R_X86_64_RELATIVE      10f0
7   000000004028     000000000008 R_X86_64_RELATIVE      4028
8   000000003fd8     000100000006 R_X86_64_GLOB_DAT 0000000000000000 _ITM_deregisterTMClone + 0
9   000000003fe0     000300000006 R_X86_64_GLOB_DAT 0000000000000000 __libc_start_main@GLIBC_2.2.5 + 0
10  000000003fe8     000400000006 R_X86_64_GLOB_DAT 0000000000000000 __gmon_start__ + 0
11  000000003ff0     000500000006 R_X86_64_GLOB_DAT 0000000000000000 _ITM_registerTMCloneTa + 0
12  000000003ff8     000600000006 R_X86_64_GLOB_DAT 0000000000000000 __cxa_finalize@GLIBC_2.2.5 + 0
13
14 Relocation section '.rela.plt' at offset 0x548 contains 1 entry:
15   Offset          Info          Type           Sym. Value      Sym. Name + Addend
16  000000004018     000200000007 R_X86_64_JUMP_SLO 0000000000000000 puts@GLIBC_2.2.5 + 0
```

ELF Sections

.dynamic section

- Contains information instructing the operating system/dynamic linker to load the binary
 - Address of important tables
 - Flags
 - Required libraries
 - Debug flags
 - INIT/FINI addresses

```
1 $ readelf --dynamic hello
2
3 Dynamic section at offset 0x2df8 contains 26 entries:
4   Tag                Type                Name/Value
5   0x0000000000000001 (NEEDED)           Shared library: [libc.so.6]
6   0x000000000000000c (INIT)             0x1000
7   0x000000000000000d (FINI)             0x11c4
8   0x0000000000000019 (INIT_ARRAY)       0x3de8
9   0x000000000000001b (INIT_ARRAYSZ)      8 (bytes)
10  0x000000000000001a (FINI_ARRAY)        0x3df0
11  0x000000000000001c (FINI_ARRAYSZ)      8 (bytes)
12  0x000000006ffffef5 (GNU_HASH)          0x308
13  0x0000000000000005 (STRTAB)            0x3d8
14  0x0000000000000006 (SYMTAB)            0x330
15  0x000000000000000a (STRSZ)             130 (bytes)
16  0x000000000000000b (SYMENT)           24 (bytes)
17  0x0000000000000015 (DEBUG)             0x0
18  0x0000000000000003 (PLTGOT)            0x4000
19  0x0000000000000002 (PLTRELSZ)          24 (bytes)
20  0x0000000000000014 (PLTREL)            RELA
21  0x0000000000000017 (JMPREL)            0x548
22  0x0000000000000007 (RELA)              0x488
23  0x0000000000000008 (RELASZ)            192 (bytes)
24  0x0000000000000009 (RELAENT)           24 (bytes)
25  0x000000006fffffff (FLAGS_1)           Flags: PIE
26  0x000000006ffffffe (VERNEED)           0x468
27  0x000000006fffffff (VERNEEDNUM)        1
28  0x000000006ffffff0 (VERSYM)            0x45a
29  0x000000006ffffff9 (RELACOUNT)         3
30  0x0000000000000000 (NULL)              0x0
```


ELF Program Headers

Overview

- Provide a **segment view** of the binary, complementing the **section view**
 - Type of segment, offset in the binary file, alignments, virtual addresses to be considered
 - **Target the operating system** that will load the program **and not the linker** as the sections do

```
1 $ readelf --wide --segments hello
2
3 Elf file type is DYN (Shared object file)
4 Entry point 0x1050
5 There are 11 program headers, starting at offset 64
6
7 Program Headers:
8   Type           Offset      VirtAddr           PhysAddr           FileSiz  MemSiz   Flg  Align
9   PHDR            0x000040    0x0000000000000040 0x0000000000000040 0x000268 0x000268 R    0x8
10  INTERP          0x0002a8    0x00000000000002a8 0x00000000000002a8 0x00001c 0x00001c R    0x1
11    [Requesting program interpreter: /lib64/ld-linux-x86-64.so.2]
12  LOAD            0x000000    0x0000000000000000 0x0000000000000000 0x000560 0x000560 R    0x1000
13  LOAD            0x001000    0x0000000000001000 0x0000000000001000 0x0001cd 0x0001cd R E  0x1000
14  LOAD            0x002000    0x0000000000002000 0x0000000000002000 0x000158 0x000158 R    0x1000
15  LOAD            0x002de8    0x0000000000003de8 0x0000000000003de8 0x000248 0x000250 RW  0x1000
16  DYNAMIC          0x002df8    0x0000000000003df8 0x0000000000003df8 0x0001e0 0x0001e0 RW  0x8
17  NOTE            0x0002c4    0x00000000000002c4 0x00000000000002c4 0x000044 0x000044 R    0x4
18  GNU_EH_FRAME     0x002010    0x0000000000002010 0x0000000000002010 0x00003c 0x00003c R    0x4
19  GNU_STACK        0x000000    0x0000000000000000 0x0000000000000000 0x000000 0x000000 RW  0x10
20  GNU_RELRO        0x002de8    0x0000000000003de8 0x0000000000003de8 0x000218 0x000218 R    0x1
```

ELF Program Headers

Types

- **LOAD**: Segment should be loaded in memory
- **INTERP**: Segment containing the name of the interpreter to be used
- **DYNAMIC**: Segment containing the **.dynamic** section, to be used by the interpreter

```
1  $ readelf --wide --segments hello
2
3  Elf file type is DYN (Shared object file)
4  Entry point 0x1050
5  There are 11 program headers, starting at offset 64
6
7  Program Headers:
8      Type           Offset      VirtAddr           PhysAddr           FileSiz  MemSiz   Flg  Align
9      PHDR            0x000040    0x0000000000000040  0x0000000000000040  0x000268 0x000268 R    0x8
10     INTERP          0x0002a8    0x00000000000002a8  0x00000000000002a8  0x00001c 0x00001c R    0x1
11         [Requesting program interpreter: /lib64/ld-linux-x86-64.so.2]
12     LOAD            0x000000    0x0000000000000000  0x0000000000000000  0x000560 0x000560 R    0x1000
13     LOAD            0x001000    0x0000000000001000  0x0000000000001000  0x0001cd 0x0001cd R E  0x1000
14     LOAD            0x002000    0x0000000000002000  0x0000000000002000  0x000158 0x000158 R    0x1000
15     LOAD            0x002de8    0x0000000000003de8  0x0000000000003de8  0x000248 0x000250 RW  0x1000
16     DYNAMIC          0x002df8    0x0000000000003df8  0x0000000000003df8  0x0001e0 0x0001e0 RW   0x8
17     NOTE            0x0002c4    0x00000000000002c4  0x00000000000002c4  0x000044 0x000044 R    0x4
18     GNU_EH_FRAME     0x002010    0x0000000000002010  0x0000000000002010  0x00003c 0x00003c R    0x4
19     GNU_STACK        0x000000    0x0000000000000000  0x0000000000000000  0x000000 0x000000 RW   0x10
20     GNU_RELRO        0x002de8    0x0000000000003de8  0x0000000000003de8  0x000218 0x000218 R    0x1
```

Dynamic Linker



Dynamic Linker

- The dynamic linker is vital for the loading process, and can aid reversing a program
 - Provide information about the loaded libraries
 - Help debugging the linking process
 - Force linking
 - And many other
- Communication is achieved through Environmental Variables
 - In the format LD_*
 - Setting a variable, or setting a variable with a specific value, activates Linker features

Dynamic Linker

LD_LIBRARY_PATH

- A list of directories in which to search for ELF libraries at execution time.
 - The items in the list are separated by either colons or semicolons
 - A zero-length directory name indicates the current working directory.
- Activating: `LD_LIBRARY_PATH=libs ./programe`
 - Will make the linker look into `./libs` while loading libraries for the program
 - Allows having a different set of libraries for the program (E.g., debug versions)

Dynamic Linker

LD_BIND_NOW

- Causes the dynamic linker to **resolve all symbols at program startup** instead of deferring function call resolution to the point when they are first referenced.
 - Especially useful for debug as all symbols point to their correct location
- Activated by setting the variable: `LD_BIND_NOW=1 progname`

```
gef> got

GOT protection: Partial RelRO | GOT functions: 4

[0x8004018] pthread_create@GLIBC_2.2.5 → 0x8001036
[0x8004020] printf@GLIBC_2.2.5 → 0x7ffffff618560
[0x8004028] pthread_exit@GLIBC_2.2.5 → 0x8001056
[0x8004030] exit@GLIBC_2.2.5 → 0x8001066
```

LD_BIND_NOW not set

```
gef> got

GOT protection: Partial RelRO | GOT functions: 4

[0x8004018] pthread_create@GLIBC_2.2.5 → 0x7ffffff797280
[0x8004020] printf@GLIBC_2.2.5 → 0x7ffffff618560
[0x8004028] pthread_exit@GLIBC_2.2.5 → 0x7ffffff7981d0
[0x8004030] exit@GLIBC_2.2.5 → 0x7ffffff5f9ea0
```

LD_BIND_NOW is set

Dynamic Linker

LD_DEBUG

- Output verbose debugging information about the operation of the dynamic linker.
 - Allows tracing the operation of the linker
 - Debug where libraries are loading from
 - Determine if libraries are being loaded and which symbols trigger the event
 - Determine the search path used looking for libraries
- The content of this variable is one of more of the following categories, separated by colons/commas, spaces:
 - help, all, bindings, files, reloc, scopes, statistics, symbols, unused, version
- Use: `LD_DEBUG=option programname`

Dynamic Linker

LD_DEBUG

```
1 $ LD_DEBUG=all ./hello_thread
2 ▼ ...
3
4 7441: relocation processing: /lib/x86_64-linux-gnu/libc.so.6 (lazy)
5 7441: symbol=_res; lookup in file=./hello_thread [0]
6 7441: symbol=_res; lookup in file=/lib/x86_64-linux-gnu/libpthread.so.0 [0]
7 7441: symbol=_res; lookup in file=/lib/x86_64-linux-gnu/libc.so.6 [0]
8 7441: binding file /lib/x86_64-linux-gnu/libc.so.6 [0] to /lib/x86_64-linux-gnu/libc.so.6 [0]: normal symbol `_res' [GLIBC_2.2.5]
9 7441: symbol=stderr; lookup in file=./hello_thread [0]
10 7441: symbol=stderr; lookup in file=/lib/x86_64-linux-gnu/libpthread.so.0 [0]
11 7441: symbol=stderr; lookup in file=/lib/x86_64-linux-gnu/libc.so.6 [0]
12 7441: binding file /lib/x86_64-linux-gnu/libc.so.6 [0] to /lib/x86_64-linux-gnu/libc.so.6 [0]: normal symbol `stderr' [GLIBC_2.2.5]
13 7441: symbol=error_one_per_line; lookup in file=./hello_thread [0]
14 7441: symbol=error_one_per_line; lookup in file=/lib/x86_64-linux-gnu/libpthread.so.0 [0]
15 7441: symbol=error_one_per_line; lookup in file=/lib/x86_64-linux-gnu/libc.so.6 [0]
16 7441: binding file /lib/x86_64-linux-gnu/libc.so.6 [0] to /lib/x86_64-linux-gnu/libc.so.6 [0]: normal symbol `error_one_per_line' [GLIBC_2.2.5]
17 7441: symbol=__morecore; lookup in file=./hello_thread [0]
18 7441: symbol=__morecore; lookup in file=/lib/x86_64-linux-gnu/libpthread.so.0 [0]
19 7441: symbol=__morecore; lookup in file=/lib/x86_64-linux-gnu/libc.so.6 [0]
20 7441: binding file /lib/x86_64-linux-gnu/libc.so.6 [0] to /lib/x86_64-linux-gnu/libc.so.6 [0]: normal symbol `__morecore' [GLIBC_2.2.5]
21 7441: symbol=__key_encryptsession_pk_LOCAL; lookup in file=./hello_thread [0]
22 7441: symbol=__key_encryptsession_pk_LOCAL; lookup in file=/lib/x86_64-linux-gnu/libpthread.so.0 [0]
23 7441: symbol=__key_encryptsession_pk_LOCAL; lookup in file=/lib/x86_64-linux-gnu/libc.so.6 [0]
24 7441: binding file /lib/x86_64-linux-gnu/libc.so.6 [0] to /lib/x86_64-linux-gnu/libc.so.6 [0]: normal symbol `__key_encryptsession_pk_LOCAL' [GLIBC_2.2.5]
25 7441: symbol=__libpthread_freeres; lookup in file=./hello_thread [0]
26 7441: symbol=__libpthread_freeres; lookup in file=/lib/x86_64-linux-gnu/libpthread.so.0 [0]
27 7441: binding file /lib/x86_64-linux-gnu/libc.so.6 [0] to /lib/x86_64-linux-gnu/libpthread.so.0 [0]: normal symbol `__libpthread_freeres'
28 7441: symbol=__progname_full; lookup in file=./hello_thread [0]
29 7441: symbol=__progname_full; lookup in file=/lib/x86_64-linux-gnu/libpthread.so.0 [0]
30 7441: symbol=__progname_full; lookup in file=/lib/x86_64-linux-gnu/libc.so.6 [0]
31 7441: binding file /lib/x86_64-linux-gnu/libc.so.6 [0] to /lib/x86_64-linux-gnu/libc.so.6 [0]: normal symbol `__progname_full' [GLIBC_2.2.5]
32 7441: symbol=__ctype32_tolower; lookup in file=./hello_thread [0]
33 7441: symbol=__ctype32_tolower; lookup in file=/lib/x86_64-linux-gnu/libpthread.so.0 [0]
34 7441: symbol=__ctype32_tolower; lookup in file=/lib/x86_64-linux-gnu/libc.so.6 [0]
```


Dynamic Linker

LD_PRELOAD

- A list of additional, user-specified, ELF shared objects to be loaded before all others.
 - This feature can be used to **selectively override functions** in other shared objects.
 - Symbols present in the provided ELF Shared objects are used instead of the original
 - Only the functions available in the shared object will be over written
- Use: `LD_PRELOAD=./liboverride.so progname`
 - Useful to provide custom implementations of any function in the program
 - Custom implementation can call the original implementation through manual symbol loading

hello_thread.c

```
1  #include <pthread.h>
2  #include <stdio.h>
3  #include <stdlib.h>
4  #define NUM_THREADS      5
5
6  void *PrintHello(void *threadid)
7  {
8      long tid;
9      tid = (long)threadid;
10     printf("Hello World! It's me, thread #%ld!\n", tid);
11     pthread_exit(NULL);
12 }
13
14 int main(int argc, char *argv[])
15 {
16     pthread_t threads[NUM_THREADS];
17     int rc;
18     long t;
19     for(t=0;t<NUM_THREADS;t++){
20         printf("In main: creating thread %ld\n", t);
21         rc = pthread_create(&threads[t], NULL, PrintHello, (void *)t);
22         if (rc){
23             printf("ERROR; return code from pthread_create() is %d\n", rc);
24             exit(-1);
25         }
26     }
27
28     pthread_exit(NULL);
29 }
```

hello_thread.c

Dynamic symbols

```
1 $ readelf --dyn-syms hello_thread
2
3 Symbol table '.dynsym' contains 10 entries:
4   Num:      Value              Size Type   Bind   Vis      Ndx Name
5       0: 0000000000000000      0 NOTYPE LOCAL DEFAULT UND
6       1: 0000000000000000      0 FUNC   GLOBAL DEFAULT UND pthread_create@GLIBC_2.2.5 (2)
7       2: 0000000000000000      0 NOTYPE WEAK   DEFAULT UND _ITM_deregisterTMCloneTab
8       3: 0000000000000000      0 FUNC   GLOBAL DEFAULT UND printf@GLIBC_2.2.5 (3)
9       4: 0000000000000000      0 FUNC   GLOBAL DEFAULT UND __libc_start_main@GLIBC_2.2.5 (3)
10      5: 0000000000000000      0 NOTYPE WEAK   DEFAULT UND __gmon_start__
11      6: 0000000000000000      0 FUNC   GLOBAL DEFAULT UND pthread_exit@GLIBC_2.2.5 (2)
12      7: 0000000000000000      0 FUNC   GLOBAL DEFAULT UND exit@GLIBC_2.2.5 (3)
13      8: 0000000000000000      0 NOTYPE WEAK   DEFAULT UND _ITM_registerTMCloneTable
14      9: 0000000000000000      0 FUNC   WEAK   DEFAULT UND __cxa_finalize@GLIBC_2.2.5 (3)
```

hello_thread.c

liboverride.c – compile with `gcc -shared -fPIC -o liboverride.so liboverride.c -ldl`

```
1  #define _GNU_SOURCE
2
3  #include <stdio.h>
4  #include <stdlib.h>
5  #include <dlfcn.h>
6  #include <unistd.h>
7  #include <sys/types.h>
8
9  void pthread_exit(){
10     void (*orig_pthread_exit)(void) = dlsym(RTLD_NEXT, "pthread_exit");
11
12     printf("pthread_exit entry\n");
13     orig_pthread_exit();
14     printf("pthread_exit exit\n");
15 }
16
17 int pthread_create(void* a, void* b, void * c, void* d){
18     int (*orig_pthread_create)(void*, void*, void*, void*) = dlsym(RTLD_NEXT, "pthread_create");
19     printf("pthread_create entry: %p %p %p %p\n", a, b, c, d);
20     int r = orig_pthread_create(a, b, c, d);
21     printf("pthread_create exit: ret=%d", r);
22     return r;
23 }
```

Manually load original function

Call original function

hello_thread.c

Left: standard execution, right: LD_PRELOAD overriding some functions

```
1 $ ./hello_thread
2 In main: creating thread 0
3 In main: creating thread 1
4 Hello World! It's me, thread #0!
5 In main: creating thread 2
6 Hello World! It's me, thread #1!
7 In main: creating thread 3
8 Hello World! It's me, thread #2!
9 In main: creating thread 4
10 Hello World! It's me, thread #3!
11 Hello World! It's me, thread #4!
```

```
1 LD_PRELOAD=./liboverride.so ./hello_thread
2 In main: creating thread 0
3 pthread_create entry: 0x7ffff9f3b5b0 (nil) 0x7f861ef96165 (nil)
4 pthread_create exit: ret=0
5 In main: creating thread 1
6 pthread_create entry: 0x7ffff9f3b5b8 (nil) 0x7f861ef96165 0x1
7 Hello World! It's me, thread #0!
8 pthread_create exit: ret=0
9 In main: creating thread 2
10 pthread_create entry: 0x7ffff9f3b5c0 (nil) 0x7f861ef96165 0x2
11 Hello World! It's me, thread #1!
12 pthread_exit entry
13 pthread_create exit: ret=0
14 In main: creating thread 3
15 pthread_exit entry
16 Hello World! It's me, thread #2!
17 pthread_exit entry
18 pthread_create entry: 0x7ffff9f3b5c8 (nil) 0x7f861ef96165 0x3
19 pthread_create exit: ret=0
20 In main: creating thread 4
21 pthread_create entry: 0x7ffff9f3b5d0 (nil) 0x7f861ef96165 0x4
22 Hello World! It's me, thread #3!
23 pthread_exit entry
24 Hello World! It's me, thread #4!
25 pthread_exit entry
26 pthread_create exit: ret=0
27 pthread_exit entry
```